

Credit Classification: Supervised Machine Learning

Gabriel D'assumpção de Carvalho

2025-07-14

Libraries

```
# install.packages("ggplot2")
# install.packages("plotly")
# install.packages("caTools")
# install.packages("e1071")
# install.packages("class")

# Graphics
library(plotly)
library(ggplot2)
library(rpart.plot)

# Data Manipulation
library(caTools)
library(dplyr)

# Machine Learning
library(e1071) # For SVM
library(class) # For KNN
library(rpart) # For Decision Trees
library(caret) # For confusionMatrix

# Warning and message suppression
options(warn = -1)
suppressMessages(library(ggplot2))
suppressMessages(library(dplyr))
suppressMessages(library(caret))

set.seed(42)
```

Introduction

Exploratory Data Analysis (EDA)

Data Import

```
# Load the dataset
df <- read.csv("/home/gabrieldadcarvalho/github/actuarial_seminar/data/original.csv")[, -1]
print(head(df))
```

##	income	age	loan	default
## 1	66155.93	59.01702	8106.5321	0
## 2	34415.15	48.11715	6564.7450	0
## 3	57317.17	63.10805	8020.9533	0
## 4	42709.53	45.75197	6103.6423	0
## 5	66952.69	18.58434	8770.0992	1
## 6	24904.06	57.47161	15.4986	0

Data Preprocessing

```
# Check statistics of the dataset
summary(df)
```

```
##      income      age      loan      default
##  Min.   :20014  Min.   :-52.42  Min.    :   1.378  Min.   :0.0000
## 1st Qu.:32796  1st Qu.: 28.99  1st Qu.: 1939.709  1st Qu.:0.0000
## Median :45789  Median : 41.32  Median : 3974.719  Median :0.0000
## Mean   :45332  Mean   : 40.81  Mean   : 4444.370  Mean   :0.1415
## 3rd Qu.:57791  3rd Qu.: 52.59  3rd Qu.: 6432.411  3rd Qu.:0.0000
## Max.   :69996  Max.   : 63.97  Max.   :13766.051  Max.   :1.0000
##                NA's      :3
```

```
print(df$age[df$age <= 0])
```

```
## [1] -28.21836 -52.42328 -36.49698      NA      NA      NA
```

Analyzing the statistics of the variables, we can see that the `age` variable has some tree NaN and negative values. Below i will converter the negative values to positive values, and for the NaN values, we discussed some imputation methods to handle these missing values.

```
# Convert age for positive values
df$age <- abs(df$age)
print(summary(df$age))
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.    NA's
##      18.06  29.03   41.35   40.92  52.59   63.97      3
```

We can see that the variable `age` has a tree NaN values. We can apply some imputation methods to handle these missing values. For example:

- Mean imputation: Replace NaN values with the mean of the column.
- Median imputation: Replace NaN values with the median of the column.
- Linear Regression imputation: Use linear regression to predict missing values based on other variables.
- Regression imputation: Use regression models to predict missing values based on other variables.
- Interpolation: Use interpolation methods to estimate missing values based on surrounding data points.
- Exploratory Data Analysis (EDA): Analyze the data to understand the distribution and value intervals of the variables, and then apply one statistics to replace the NaN values.

Exploratory Data Analysis (EDA)

```
p1 <- plot_ly(df, x = ~age, type = "histogram", name = "Age") %>%
  layout(title = "Age", xaxis = list(title = "Age"), yaxis = list(title = "Count"), showlegend = FALSE)

p2 <- plot_ly(df, x = ~income, type = "histogram", name = "Income") %>%
  layout(title = "Income", xaxis = list(title = "Income"), yaxis = list(title = "Count"), showlegend = FALSE)

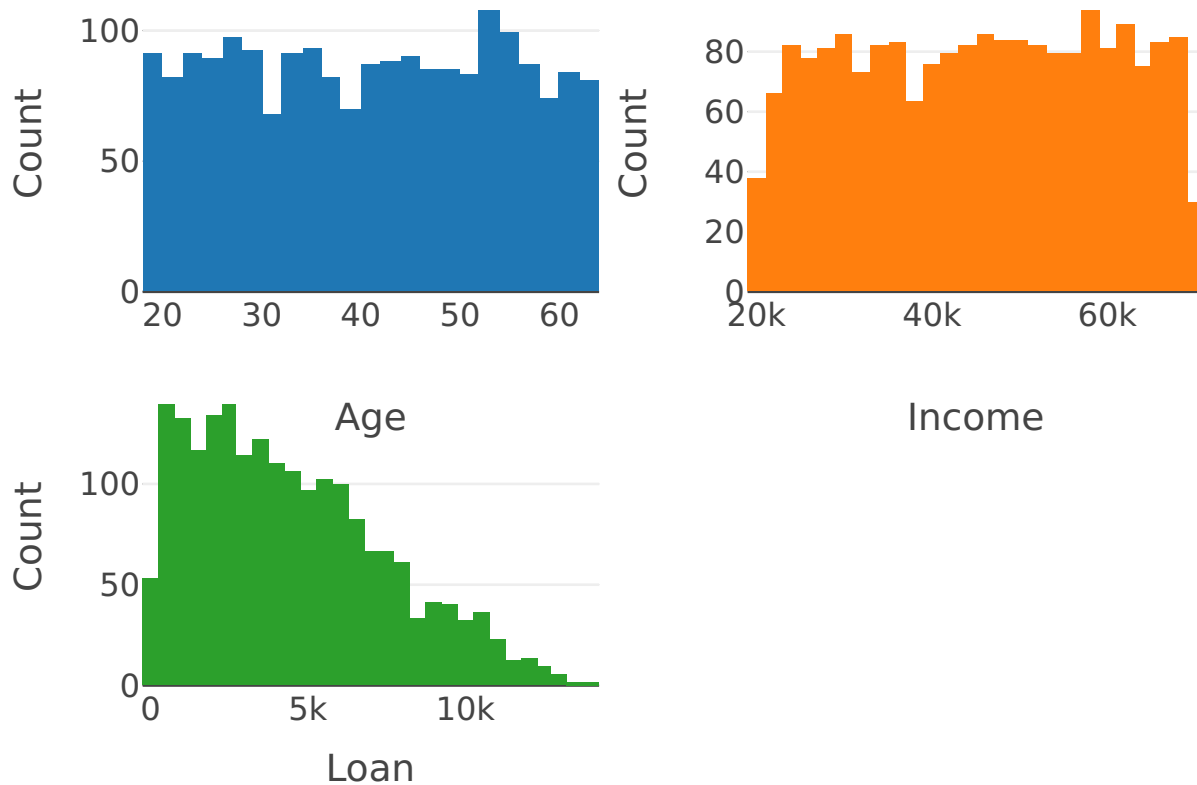
p3 <- plot_ly(df, x = ~loan, type = "histogram", name = "Loan") %>%
  layout(title = "Loan", xaxis = list(title = "Loan"), yaxis = list(title = "Count"), showlegend = FALSE)

subplot(p1, p2, p3,
  nrows = 2, margin = 0.07,
  titleX = TRUE, titleY = TRUE,
  shareX = FALSE, shareY = FALSE)
```

```
) %>%
  layout(title = "Distribution of Explanatory Variables")
```

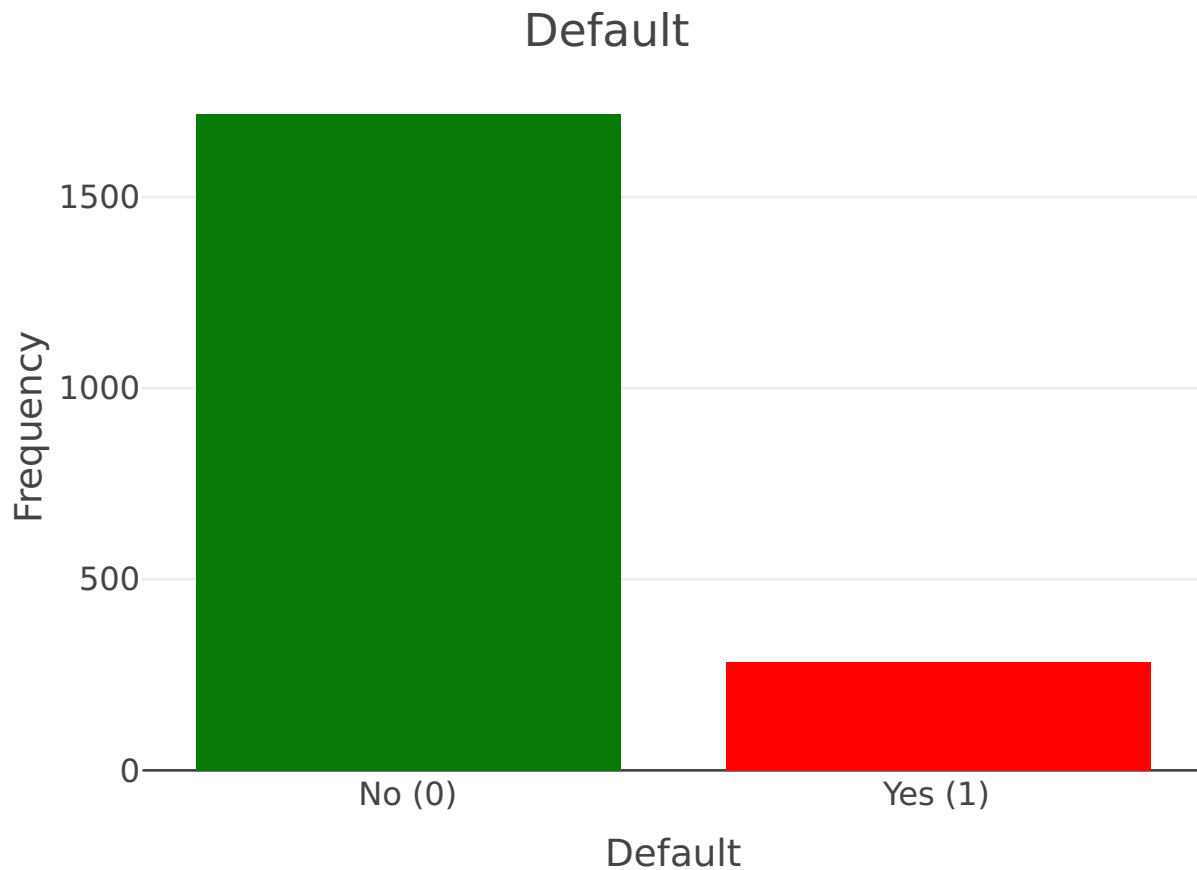
```
## `google-chrome` and `chromium-browser` were not found. Try setting the `CHROMOTE_CHROME` environment
```

Distribution of Explanatory Variables



The `age` and `income` variables have a similar uniform distribution, while the `loan` variable has an asymmetric positive distribution. This affirmation is confirmed by the summary statistics of the dataset and the above plots.

```
df %>%
  count(default) %>%
  plot_ly(
    x = ~ factor(default, levels = c(0, 1), labels = c("No (0)", "Yes (1)")),
    y = ~n,
    type = "bar",
    color = ~ factor(default),
    colors = c("#067906", "red"),
    name = "Default Variable"
  ) %>%
  layout(
    title = "Default",
    xaxis = list(title = "Default", type = "category"),
    yaxis = list(title = "Frequency"),
    showlegend = FALSE
  )
```



```
corrPearson <- cor(df[, !names(df) %in% "default"], method = "pearson", use = "pairwise.complete.obs")
corrSpearman <- cor(df[, !names(df) %in% "default"], method = "spearman", use = "pairwise.complete.obs")
print(corrPearson)
```

```
##           income           age           loan
## income  1.00000000 -0.034000535 0.441116504
## age     -0.03400054  1.000000000 0.006440323
## loan     0.44111650  0.006440323 1.000000000
```

```
print(corrSpearman)
```

```
##           income           age           loan
## income  1.00000000 -0.034486684 0.401601061
## age     -0.03448668  1.000000000 0.009835956
## loan     0.40160106  0.009835956 1.000000000
```

Linear Regression For Imputation data

For imputation, will be used the Linear Regression algorithm, which is a simple and effective methods for handling missing values. The Linear Regression algorithm works by finding the relationship between the target variable and the other features to predict the missing values.

```
# Get the missing values in the 'age' column
ageNan <- df[is.na(df$age), ]
print(ageNan)
```

```
##      income age      loan default
## 29 59417.81  NA 2082.626         0
```

```
## 31 48528.85 NA 6155.785 0
## 32 23526.30 NA 2862.010 0
```

```
# Remove rows with NaN in 'age' for training and testing
dfNN <- df[!is.na(df$age), ]
print(summary(dfNN))
```

```
##      income      age      loan      default
## Min.   :20014 Min.   :18.06 Min.   :  1.378 Min.   :0.0000
## 1st Qu.:32805 1st Qu.:29.03 1st Qu.: 1936.813 1st Qu.:0.0000
## Median :45789 Median :41.35 Median : 3977.287 Median :0.0000
## Mean   :45334 Mean   :40.92 Mean   : 4445.488 Mean   :0.1417
## 3rd Qu.:57788 3rd Qu.:52.59 3rd Qu.: 6440.861 3rd Qu.:0.0000
## Max.   :69996 Max.   :63.97 Max.   :13766.051 Max.   :1.0000
```

Split the dataset into training and testing sets

```
# Split the dataset into training and testing sets
split <- sample.split(dfNN$age, SplitRatio = 0.8)
train <- subset(dfNN, split == TRUE)
test <- subset(dfNN, split == FALSE)
```

Normalize the numeric columns

```
# Select numeric columns, excluding 'default'
numeric_cols <- sapply(train, is.numeric)
cols_for_stats <- names(train)[numeric_cols & names(train) != "default"]

# Calculate statistics only for the selected columns
means <- colMeans(train[, cols_for_stats], na.rm = TRUE)
sds <- apply(train[, cols_for_stats], 2, sd, na.rm = TRUE)

# Z-score normalization
for (c in colnames(train)[numeric_cols & names(train) != "default"]){
  if (is.numeric(train[[c]])) {
    train[[c]] <- (train[[c]] - means[c]) / sds[c]
    test[[c]] <- (test[[c]] - means[c]) / sds[c]
    ageNan[[c]] <- (ageNan[[c]] - means[c]) / sds[c]
    df[[c]] <- (df[[c]] - means[c]) / sds[c]
  }
}

print(summary(train))
```

```
##      income      age      loan      default
## Min.   :-1.76492 Min.   :-1.73218 Min.   :-1.4454 Min.   :0.0000
## 1st Qu.: -0.87437 1st Qu.: -0.89946 1st Qu.: -0.8240 1st Qu.:0.0000
## Median : 0.03517 Median : 0.03068 Median :-0.1510 Median :0.0000
## Mean   : 0.00000 Mean   : 0.00000 Mean   : 0.0000 Mean   :0.1403
## 3rd Qu.: 0.87019 3rd Qu.: 0.89202 3rd Qu.: 0.6672 3rd Qu.:0.0000
## Max.   : 1.72498 Max.   : 1.71814 Max.   : 3.0545 Max.   :1.0000
```

```
print(summary(test))
```

```
##      income      age      loan      default
```

```
## Min.      :-1.768317   Min.      :-1.71332   Min.      :-1.4234   Min.      :0.0000
## 1st Qu.: -0.869711   1st Qu.: -0.94137   1st Qu.: -0.7226   1st Qu.: 0.0000
## Median :  0.030767   Median : -0.03091   Median : -0.1195   Median : 0.0000
## Mean    :  0.006546   Mean     :-0.06830   Mean     : 0.0373   Mean     : 0.1475
## 3rd Qu.:  0.877523   3rd Qu.:  0.74539   3rd Qu.:  0.6149   3rd Qu.: 0.0000
## Max.    :  1.703976   Max.     :  1.71162   Max.     :  2.9491   Max.     : 1.0000
```

```
print(summary(ageNan))
```

```
##      income      age      loan      default
## Min.      :-1.5229   Min.      : NA   Min.      :-0.7650   Min.      :0
## 1st Qu.: -0.6491   1st Qu.:  NA   1st Qu.: -0.6376   1st Qu.:0
## Median :  0.2246   Median :  NA   Median : -0.5102   Median :0
## Mean     :-0.1042   Mean     :NaN   Mean     :-0.2362   Mean     :0
## 3rd Qu.:  0.6051   3rd Qu.:  NA   3rd Qu.:  0.0282   3rd Qu.:0
## Max.     :  0.9857   Max.     :  NA   Max.     :  0.5666   Max.     :0
##                                     NA's      :3
```

Train the Linear Regression model

```
# Train a linear regression model to predict 'age'
model <- lm(age ~ ., data = train)
```

```
print(summary(model))
```

```
##
## Call:
## lm(formula = age ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.02876 -0.62141  0.04569  0.63442  2.07822
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.22153    0.02372   9.338 < 2e-16 ***
## income      -0.18383    0.02448  -7.509 9.88e-14 ***
## loan         0.29675    0.02654  11.180 < 2e-16 ***
## default     -1.57939    0.06889 -22.927 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8658 on 1593 degrees of freedom
## Multiple R-squared:  0.2518, Adjusted R-squared:  0.2504
## F-statistic: 178.7 on 3 and 1593 DF, p-value: < 2.2e-16
```

Test the model

```
# Predict missing 'age' values in the test set
predictedAge <- predict(model, newdata = test)
```

```
# Mean Squared Error (MSE) for the predictions
mse <- 0
for (i in 1:length(predictedAge)) {
  mse = mse + (predictedAge[i] - test$age[i])^2
}
```

```
}

mse = mse / length(predictedAge)
print(mse)
```

```
##          1
## 0.7245957
```

Impute missing values in the original dataset

```
# Impute missing values in the original dataset
ageNan$age <- predict(model, newdata = ageNan)

# Replace NaN values in the original dataset with the predicted values
df$age[is.na(df$age)] <- ageNan$age

# Reverse Z-score normalization for the imputed values
for (c in colnames(df[(numeric_cols & names(df) != "default")])) {
  if (is.numeric(df[[c]])) {
    df[[c]] <- (df[[c]] * sds[c]) + means[c]
    ageNan[[c]] <- (ageNan[[c]] * sds[c]) + means[c]
  }
}

print(head(ageNan))
```

```
##      income      age      loan default
## 29 59417.81 38.62265 2082.626        0
## 31 48528.85 45.74307 6155.785        0
## 32 23526.30 45.76578 2862.010        0
```

```
print(summary(df))
```

```
##      income      age      loan      default
## Min.   :20014   Min.   :18.06   Min.    :  1.378   Min.    :0.0000
## 1st Qu.:32796   1st Qu.:29.06   1st Qu.: 1939.709   1st Qu.:0.0000
## Median :45789   Median :41.38   Median : 3974.719   Median :0.0000
## Mean   :45332   Mean   :40.93   Mean    : 4444.370   Mean    :0.1415
## 3rd Qu.:57791   3rd Qu.:52.58   3rd Qu.: 6432.411   3rd Qu.:0.0000
## Max.   :69996   Max.    :63.97   Max.    :13766.051   Max.    :1.0000
```

Predict Default

Data Preparation

```
split <- sample.split(df$default, SplitRatio = 0.8)
train <- subset(df, split == TRUE)
test  <- subset(df, split == FALSE)

means <- colMeans(train[, cols_for_stats], na.rm = TRUE)
sds <- apply(train[, cols_for_stats], 2, sd, na.rm = TRUE)

# Z-score normalization
for (c in colnames(train[(numeric_cols & names(train) != "default")])) {
```



```

if (is.numeric(train[[c]])) {
  train[[c]] <- (train[[c]] - means[c]) / sds[c]
  test[[c]] <- (test[[c]] - means[c]) / sds[c]
}
}

```

Logistic Regression

```

# Train a logistic regression model
logistic_model <- glm(default ~ ., data = train, family = binomial(link = "logit"))

print(summary(logistic_model))

```

```

##
## Call:
## glm(formula = default ~ ., family = binomial(link = "logit"),
##      data = train)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -8.1968      0.6598  -12.42  <2e-16 ***
## income       -3.6457      0.3643  -10.01  <2e-16 ***
## age          -5.0026      0.4248  -11.78  <2e-16 ***
## loan         5.6158      0.4907   11.45  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1303.12  on 1599  degrees of freedom
## Residual deviance:  323.71  on 1596  degrees of freedom
## AIC: 331.71
##
## Number of Fisher Scoring iterations: 9

```

```

# Predict on the test set
logistic_pred <- predict(logistic_model, newdata = test, type = "response")
logistic_pred_class <- ifelse(logistic_pred > 0.5, 1, 0)

head(data.frame(Prediction = logistic_pred_class, Probability = logistic_pred))

```

```

##      Prediction  Probability
## 3              0 2.330707e-06
## 4              0 1.948964e-03
## 5              1 9.395945e-01
## 9              0 7.029250e-06
## 10             0 2.947396e-04
## 13             0 1.853924e-02

```

```

# Confusion matrix for logistic regression
confusionMatrix(as.factor(test$default), as.factor(logistic_pred_class), dnn = c("Reference", "Prediction"))

```

```

## Confusion Matrix and Statistics
##
##              Prediction

```

```
## Reference    0    1
##           0 327  16
##           1  14  43
##
##           Accuracy : 0.925
##           95% CI : (0.8947, 0.9488)
##       No Information Rate : 0.8525
##       P-Value [Acc > NIR] : 6.876e-06
##
##           Kappa : 0.6975
##
##  McNemar's Test P-Value : 0.8551
##
##           Sensitivity : 0.9589
##           Specificity : 0.7288
##       Pos Pred Value : 0.9534
##       Neg Pred Value : 0.7544
##           Prevalence : 0.8525
##       Detection Rate : 0.8175
##       Detection Prevalence : 0.8575
##       Balanced Accuracy : 0.8439
##
##       'Positive' Class : 0
##
```

Support Vector Machine (SVM)

```
# Train a Support Vector Machine (SVM) model
svm_model <- svm(default ~ ., data = train, kernel = "radial", cost = 1, gamma = 3)

summary(svm_model)
```

```
##
## Call:
## svm(formula = default ~ ., data = train, kernel = "radial", cost = 1,
##      gamma = 3)
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: radial
##      cost:   1
##     gamma:   3
##   epsilon:  0.1
##
## Number of Support Vectors:  524
```

```
# Predict on the test set using SVM
svm_pred <- predict(svm_model, newdata = test)
svm_pred_class <- ifelse(svm_pred >= 0.5, 1, 0)

head(data.frame(Previsto = svm_pred_class, Prob = svm_pred))
```

```
##   Previsto      Prob
```

```
## 3      0  0.03786639
## 4      0  0.03801636
## 5      1  0.85102751
## 9      0  0.03643876
## 10     0  0.02738613
## 13     0 -0.01139708
```

```
levels_ref <- c("0", "1")
```

```
reference <- factor(test$default, levels = levels_ref)
prediction <- factor(svm_pred_class, levels = levels_ref)
```

```
# Confusion matrix for SVM
```

```
confusionMatrix(reference, prediction, dnn = c("Reference", "Prediction"))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##      Prediction
```

```
## Reference  0    1
```

```
##      0 342    1
```

```
##      1   9   48
```

```
##
```

```
##      Accuracy : 0.975
```

```
##      95% CI : (0.9545, 0.9879)
```

```
##      No Information Rate : 0.8775
```

```
##      P-Value [Acc > NIR] : 1.758e-12
```

```
##
```

```
##      Kappa : 0.8913
```

```
##
```

```
##      McNemar's Test P-Value : 0.02686
```

```
##
```

```
##      Sensitivity : 0.9744
```

```
##      Specificity : 0.9796
```

```
##      Pos Pred Value : 0.9971
```

```
##      Neg Pred Value : 0.8421
```

```
##      Prevalence : 0.8775
```

```
##      Detection Rate : 0.8550
```

```
##      Detection Prevalence : 0.8575
```

```
##      Balanced Accuracy : 0.9770
```

```
##
```

```
##      'Positive' Class : 0
```

```
##
```

Decision tree

Data Preparation

```
split <- sample.split(df$default, SplitRatio = 0.8)
```

```
train <- subset(df, split == TRUE)
```

```
test <- subset(df, split == FALSE)
```

```
means <- colMeans(train[, cols_for_stats], na.rm = TRUE)
```

```
sds <- apply(train[, cols_for_stats], 2, sd, na.rm = TRUE)
```

```
trainT <- train["default"]
```

```

testT <- test["default"]

trainT$ageInterval <- cut(train$age, breaks = c(18, 40, 60, Inf), right = FALSE)
testT$ageInterval <- cut(test$age, breaks = c(18, 40, 60, Inf), right = FALSE)

# Menos intervalos para income
trainT$incomeInterval <- cut(train$income, breaks = c(20000, 50000, 100000, Inf), right = FALSE)
testT$incomeInterval <- cut(test$income, breaks = c(20000, 50000, 100000, Inf), right = FALSE)

# Menos intervalos para loan
trainT$loanInterval <- cut(train$loan, breaks = c(0, 5000, 10000, Inf), right = FALSE)
testT$loanInterval <- cut(test$loan, breaks = c(0, 5000, 10000, Inf), right = FALSE)

ageClass <- table(trainT$ageInterval, trainT$default)
print(ageClass)

##
##           0    1
## [18,40)  523  226
## [40,60)  713    0
## [60,Inf)  138    0

incomeClass <- table(trainT$incomeInterval, trainT$default)
print(incomeClass)

##
##           0    1
## [2e+04,5e+04)  810  135
## [5e+04,1e+05)  564   91
## [1e+05,Inf)     0    0

loanClass <- table(trainT$loanInterval, trainT$default)
print(loanClass)

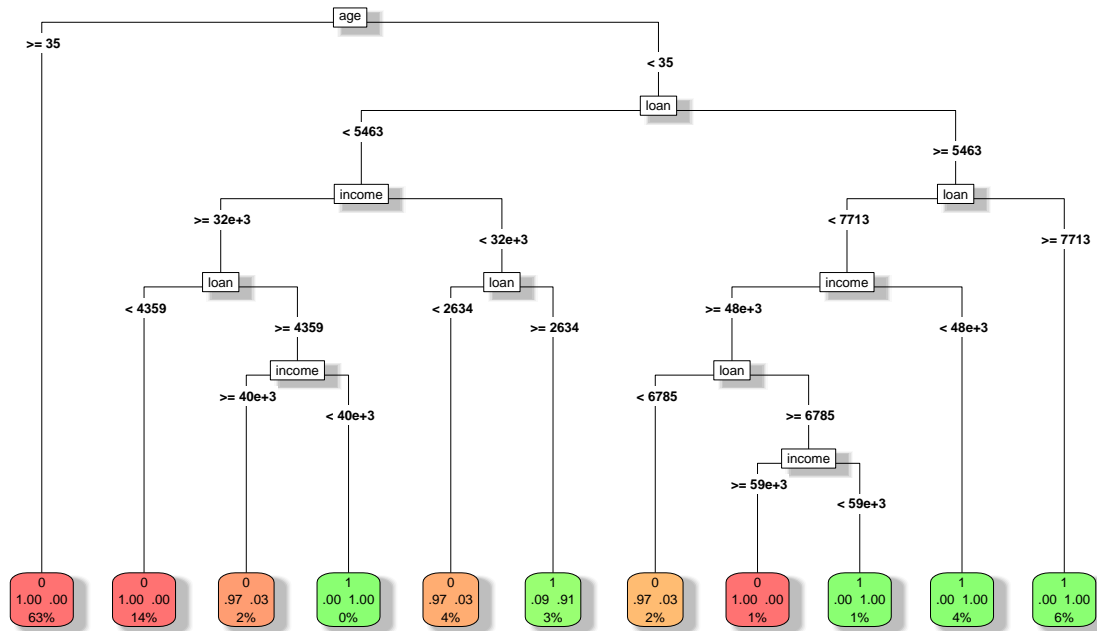
##
##           0    1
## [0,5e+03)     922   50
## [5e+03,1e+04) 392  141
## [1e+04,Inf)    60   35

dt_model <- rpart(default ~., data=train, method = "class")

rpart.plot(dt_model, type=5, extra=104, main="Decision Tree for Default Prediction", box.palette = "RdY")

```

Decision Tree for Default Prediction



Predicting

```
predictions <- predict(dt_model, newdata = test, type = "class")
```

```
# Confusion matrix for Decision Tree
```

```
confusionMatrix(as.factor(test$default), as.factor(predictions), dnn = c("Reference", "Prediction"))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##      Prediction
```

```
## Reference  0    1
```

```
##           0 340   3
```

```
##           1   5  52
```

```
##
```

```
##           Accuracy : 0.98
```

```
##           95% CI : (0.961, 0.9913)
```

```
## No Information Rate : 0.8625
```

```
## P-Value [Acc > NIR] : <2e-16
```

```
##
```

```
##           Kappa : 0.9169
```

```
##
```

```
## McNemar's Test P-Value : 0.7237
```

```
##
```

```
##           Sensitivity : 0.9855
```

```
##           Specificity : 0.9455
```

```
## Pos Pred Value : 0.9913
```

```
## Neg Pred Value : 0.9123
```

```
## Prevalence : 0.8625
```

```
## Detection Rate : 0.8500
```

```
## Detection Prevalence : 0.8575
```

```
##      Balanced Accuracy : 0.9655
##
##      'Positive' Class : 0
##
```